

Distributed DBs and ACID - Pessimistic Concurrency:

- ACID transactions
 - Focuses on “data safety”
 - considered a pessimistic concurrency model because it assumes one transaction has to protect itself from other transactions
 - IOW, it assumes that if something can go wrong, it will.
 - Conflicts are prevented by locking resources until a transaction is complete (there are both read and write locks)
 - Write Lock Analogy → borrowing a book from a library... If you have it, no one else can.

Optimistic Concurrency:

- Transactions do not obtain locks on data when they read or write
- *Optimistic* because it assumes conflicts are unlikely to occur
 - Even if there is a conflict, everything will still be OK.
- But how?
 - Add last update timestamp and version number columns to every table... read them when changing. THEN, check at the end of transaction to see if any other transaction has caused them to be modified.
- Low Conflict Systems (backups, analytical dbs, etc.)
 - Read heavy systems
 - the conflicts that arise can be handled by rolling back and re-running a transaction that notices a conflict.
 - So, optimistic concurrency works well - allows for higher concurrency
- High Conflict Systems
 - rolling back and rerunning transactions that encounter a conflict → less efficient

So, a locking scheme (pessimistic model) might be preferable

NoSQL:

- “NoSQL” first used in 1998 by Carlo Strozzi to describe his relational database system that *did not use SQL*.
- More common, modern meaning is “Not Only SQL”
- *But*, sometimes thought of as non-relational DBs
- Idea originally developed, in part, as a response to processing unstructured web-based data.

CAP Theorem Review:

You can have 2, but not 3, of the following:

- Consistency*: Every user of the DB has an identical view of the data at any given instant
 - Availability: In the event of a failure, the database system remains operational
 - Partition Tolerance: The database can maintain operations in the event of the network's failing between two segments of the distributed system
-
- Consistency + Availability: System always responds with the latest data and every request gets a response, but may not be able to deal with network partitions
 - Consistency + Partition Tolerance: If system responds with data from the distrib. system, it is always the latest, else data request is dropped.
 - Availability + Partition Tolerance: System always sends responses based on distributed store, but may not be the absolute latest data.

ACID Alternative for Distribution Systems - BASE:

- Basically Available
 - Guarantees the availability of the data (per CAP), but response can be “failure”/“unreliable” because the data is in an inconsistent or changing state
 - System appears to work most of the time
- Soft State - The state of the system could change over time, even w/o input. Changes could be result of *eventual consistency*.
 - Data stores don't have to be write-consistent
 - Replicas don't have to be mutually consistent
- Eventual Consistency - The system will eventually become consistent
 - All writes will eventually stop so all nodes/replicas can be updated

Key-Value Databases:

Key-Value Stores:

key = value

- Key-value stores are designed around:
 - *simplicity*
 - the data model is extremely simple
 - comparatively, tables in a RDBMS are very complex.
 - lends itself to simple CRUD ops and API creation
- *speed*

- usually deployed as in-memory DB
- retrieving a *value* given its *key* is typically a $O(1)$ op b/c hash tables or similar data structs used under the hood
- no concept of complex queries or joins... they slow things down
- *scalability*
 - Horizontal Scaling is simple - add more nodes
 - Typically concerned with *eventual consistency*, meaning in a distributed environment, the only guarantee is that all nodes will *eventually* converge on the same value.

KV DS Use Cases:

- EDA/Experimentation Results Store
 - store intermediate results from data preprocessing and EDA
 - store experiment or testing (A/B) results w/o prod db
- Feature Store
 - store frequently accessed feature → low-latency retrieval for model training and prediction
- Model Monitoring
 - store key metrics about performance of model, for example, in real-time inferencing.
- Storing Session Information
 - everything about the current *session* can be stored via a single PUT or POST and retrieved with a single GET VERY Fast
- User Profiles & Preferences
 - User info could be obtained with a single GET operation... language, TZ, product or UI preferences
- Shopping Cart Data
 - Cart data is tied to the user
 - needs to be available across browsers, machines, sessions
- Caching Layer:
 - In front of a disk-based database

Redis DB:

- Redis (Remote Directory Server)
 - Open source, in-memory database
 - Sometimes called a data structure store
 - Primarily a KV store, but can be used with other models: Graph, Spatial, Full Text Search, Vector, Time Series
 - It is considered an in-memory database system, but...

- Supports durability of data by: a) essentially saving snapshots to disk at specific intervals or b) append-only file which is a journal of changes that can be used for roll-forward if there is a failure
- Originally developed in 2009 in C++
- Can be very fast ... > 100,000 SET ops / second
- Rich collection of commands
- Does NOT handle complex data. No secondary indexes. Only supports lookup by Key.

Redis Data Types:

Keys:

- usually strings but can be any binary sequence


Values:

- Strings
- Lists (linked lists)
- Sets (unique unsorted string elements)
- Sorted Sets
- Hashes (string → string)
- Geospatial data

Setting Up Redis in Docker:

- In Docker Desktop, search for Redis.
- Pull/Run the latest image (see above)
 - Optional Settings: add 6379 to Ports to expose that port so we can connect to it.
- Normally, you would not expose the Redis port for security reasons
 - If you did this in a prod environment, major security hole.
 - Notice, we didn't set a password...

Connecting from DataGrip:

- File > New > Data Source > Redis
- Give the Data Source a Name
- Make sure the port is 6379
- Test the connection 

Redis Database and Interaction:

- Redis provides 16 databases by default
 - They are numbered 0 to 15
 - There is no other name associated
- Direct interaction with Redis is through a set of commands related to setting and getting k/v pairs (and variations)
- Many language libraries available as well.

Foundation Data Type - String:

- Sequence of bytes - text, serialized objects, bin arrays
- Simplest data type
- Maps a string to another string
- Use Cases:
 - caching frequently accessed HTML/CSS/JS fragments
 - config settings, user settings info, token management
 - counting web page/app screen views OR rate limiting

Some Initial Basic Commands:

- SET /path/to/resource 0
SET user:1 "John Doe"
GET /path/to/resource
EXISTS user:1
DEL user:1
KEYS user*
- SELECT 5
 - select a different database

Some Basic Commands:

SET someValue 0

INCR someValue #increment by 1

INCRBY someValue 10 #increment by 10

DECR someValue #decrement by 1

DECRBY someValue 5 #decrement by 5

- INCR parses the value as int and increments (or adds to value)
- SETNX key value

only sets value to key if key does not already exist

Hash Type:

Value of KV entry is a collection of *field-value* pairs

- Use Cases:
 - Can be used to represent basic objects/structures
 - number of field/value pairs per hash is $2^{32}-1$
 - practical limit: available system resources (e.g. memory)
 - Session information management
 - User/Event tracking (could include TTL)
 - Active Session Tracking (all sessions under one hash key)

Hash Commands:

HSET bike:1 model Demios brand Ergonom price 1971

HGET bike:1 model

HGET bike:1 price

HGETALL bike:1

HMGET bike:1 model price *weight*

HINCRBY bike:1 price 100

List Type:

- Value of KV Pair is linked lists of string values
- Use Cases:
 - implementation of stacks and queues
 - queue management & message passing queues (producer/consumer model)
 - logging systems (easy to keep in chronological order)
 - build social media streams/feeds
 - message history in a chat application
 - batch processing by queueing up a set of tasks to be executed sequentially at a later time

Linked Lists Crash Course:

- Sequential data structure of linked nodes (instead of contiguously allocated memory)
- Each node points to the next element of the list (except the last one - points to *nil/null*)
- $O(1)$ to insert new value at front or insert new value at end

List Commands - Queue:

Queue-like Ops

LPUSH bikes:repairs bike:1

LPUSH bikes:repairs bike:2

RPOP bikes:repairs

RPOP bikes:repairs

List Commands - Stack:

Stack-like Ops

LPUSH bikes:repairs bike:1

LPUSH bikes:repairs bike:2

LPOP bikes:repairs

LPOP bikes:repairs

List Commands - Others:

Other List Ops

LLEN mylist

LRANGE <key> <start> <stop>

LRANGE mylist 0 3

LRANGE mylist 0 0

LRANGE mylist -2 -1

JSON Type:

- Full support of the JSON standard
- Uses JSONPath syntax for parsing/navigating a JSON document
- Internally, stored in binary in a tree-structure → fast access to sub elements

Set Type:

- Unordered collection of unique strings (members)
- Use Cases:
 - track unique items (IP addresses visiting a site, page, screen)
 - primitive relation (set of all students in DS4300)
 - access control lists for users and permission structures
 - social network friends lists and/or group membership
- Supports set operations!!

Set Commands:

SADD ds4300 "Mark"

SADD ds4300 "Sam"

SADD cs3200 "Nick"

SADD cs3200 "Sam"

SISMEMBER ds4300 "Mark"

SISMEMBER ds4300 "Nick"

SCARD ds4300

SADD ds4300 "Mark"

SADD ds4300 "Sam"

SADD cs3200 "Nick"

SADD cs3200 "Sam"

SCARD ds4300

SINTER ds4300 cs3200

SDIFF ds4300 cs3200

SREM ds4300 "Mark"

SRANDMEMBER ds4300