

Searching:

- Searching is the most common operation performed by a database system
- In SQL, the SELECT statement is arguably the most versatile / complex.
- Baseline for efficiency is Linear Search
 - Start at the beginning of a list and proceed element by element until:
 - You find what you're looking for
 - You get to the last element and haven't found it
- Record - A collection of values for attributes of a single entity instance; a row of a table
- Collection - a set of records of the same entity type; a table
 - Trivially, stored in some sequential order like a list
- Search Key - A value for an attribute from the entity type
 - Could be ≥ 1 attribute

List of Records:

- If each record takes up x bytes of memory, then for n records, we need $n*x$ bytes of memory.
- Contiguously Allocated List
 - All $n*x$ bytes are allocated as a single "chunk" of memory
- Linked List
 - Each record needs x bytes + additional space for 1 or 2 memory addresses
 - Individual records are linked together in a type of chain using memory addresses
- Arrays are faster for random access, but slow for inserting anywhere but the end
- Linked Lists are faster for inserting anywhere in the list, but slower for random access

Observations:

- Arrays
 - fast for random access
 - slow for random insertions
- Linked Lists
 - slow for random access
 - fast for random insertions

Binary Search:

- Input: array of values in sorted order, target value
- Output: the location (index) of where target is located or some value indicating target was not found

```
def binary_search(arr, target)
    left, right = 0, len(arr) - 1
    while left <= right:
        mid = (left + right) // 2
        if arr[mid] == target:
            return mid
        elif arr[mid] < target:
            left = mid + 1
        else:
            right = mid - 1
    return -1
```

Time Complexity:

- Linear Search
 - Best case: target is found at the first element; only 1 comparison
 - Worst case: target is not in the array; n comparisons
 - Therefore, in the worst case, linear search is $O(n)$ time complexity.
- Binary Search
 - Best case: target is found at mid ; 1 comparison (inside the loop)
 - Worst case: target is not in the array; $\log_2 n$ comparisons
 - Therefore, in the worst case, binary search is $O(\log_2 n)$ time complexity.

Back to Database Searching:

- Assume data is stored on disk by column id's value
- Searching for a specific id = fast.
- But what if we want to search for a specific *specialVal*?
 - Only option is linear scan of that column
- Can't store data on disk sorted by both id and *specialVal* (at the same time)
 - data would have to be duplicated → space inefficient

We need an external data structure to support faster searching by *specialVal* than a linear scan.

- 1) An array of tuples (specialVal, rowNum) sorted by specialVal
 - a) We could use Binary Search to quickly locate a particular specialVal and find its corresponding row in the table
 - b) But, every insert into the table would be like inserting into a sorted array - slow...
- 2) A linked list of tuples (specialVal, rowNum) sorted by specialVal
 - a) searching for a specialVal would be slow - linear scan required
 - b) But inserting into the table would theoretically be quick to also add to the list.

Something with Fast Insert and Fast Search?:

Binary Search Tree - a binary tree where every node in the left subtree is less than its parent and every node in the right subtree is greater than its parent.